

Exercise 4, Part 1: Two Skills

Session: 4 — Skills, Hooks & Automation

Duration: 15 minutes. Part 2 is a separate sheet.

Project: The same BookStore API.

Goal

The commands you have used since session 2 are markdown files in this repo. Now you make two of your own: copy a ready-made `/commit` skill, then let Claude write a `/changelog` skill and write its `description` yourself.

Part 2, after the slides on hooks, adds a hook that runs `/changelog` after every commit. It has its own sheet.

Start here

- Open a terminal in the `bookstore` folder and start `claude`. Everything you create goes in `bookstore/.claude/`.
- Training mode is on. Some prompts on this sheet start with an experiment tag in square brackets. The tag switches training mode off for that one prompt. Copy those prompts exactly, tag included.
- Never edit `CLAUDE.md`. Sessions 5 to 8 need it.
- No coach this session: the *Done when* lines are your checklist, and `/verify-exercise 4` grades your files. Stuck? Type *just tell me*.

Tasks

1. Copy the `/commit` skill and run it

This task produces a branch, the file `.claude/skills/commit/SKILL.md`, and one commit.

1. **Create the branch.** Sessions 5 to 8 reuse this code. You delete the branch after the closing round.

```
git switch -c session4-playground
```

2. **Let Claude copy the skill.** The skill is ready in `exercises/starters/commit-skill.md`. Paste:

[Exercise 4 experiment – execute directly, no leading questions.] Copy exercises/starters/commit-skill.md to .claude/skills/commit/SKILL.md. Change nothing in the content.

3. **Make a small change by hand.** Open any file in bookstore , add one comment line, save it. Do not ask Claude.
4. **Run the skill.** Type /commit and watch what it does.
5. **Read three lines of the skill.** Open .claude/skills/commit/SKILL.md in your editor. Three lines in it matter most. "Execute directly — no leading questions" switches training mode off inside the skill; without it, Claude asks a question instead of committing. disable-model-invocation: true means only typing /commit runs it. Imperative mood means the message starts with a present-tense verb: "Add", "Fix", "Remove". Task 2's changelog is built from those messages.

Done when: the commit exists, its message starts with a present-tense verb, and nothing was pushed.

When the trainer calls task 2, start it, even if /commit has not run yet.

2. Let Claude write /changelog , then write its description

This task produces .claude/skills/changelog/SKILL.md and, next to it, common-changelog-spec.md .

1. **One prompt for both files.** A skill can carry extra files, loaded only when it runs (*Anatomy of a Skill*). So the specification lives next to the skill, not inside it. Paste:

[Exercise 4 experiment – execute directly, no leading questions.] Read <https://common-changelog.org/> and save the whole specification as markdown to `.claude/skills/changelog/common-changelog-spec.md`. Keep the headings and the examples. Leave out the website menu and footer. Then create `.claude/skills/changelog/SKILL.md`. Frontmatter: the line ``name: changelog`` and the line ``description: (I write this)``. No other frontmatter lines. Body: a numbered list of instructions that make you do all of the following.

- Read `CHANGELOG.md` in the project root. If the file does not exist, create it with the heading ``# Changelog``.
- Read `common-changelog-spec.md`, next to this file, before changing the changelog.
- Look at the recent commits and at the existing version tags.
- Add only new commits. Entries that are already in the changelog stay as they are. New entries go in a ``## Unreleased`` section at the top.
- Sort entries under ``### Changed``, ``### Added``, ``### Removed`` and ``### Fixed``, in that order.
- Write each entry as one sentence that starts with a present-tense verb. End it with the short commit hash in round brackets, like `(a1b2c3d)`.
- A breaking change forces other people to change their own code. Put `**Breaking:**` in front of such an entry, and put those entries first inside their group.
- Leave out changes a reader does not care about: files whose name starts with a dot, formatting-only changes, and developer tools.

If your network blocks the page, ask Claude to write the spec file from what it knows about Common Changelog.

2. **Check two things.** `common-changelog-spec.md` holds the group names and the rules, not an empty page. The `SKILL.md` frontmatter has **no** `disable-model-invocation` line: Claude must be able to start this skill by itself. The hook in Part 2 only suggests it; Claude decides.
3. **Write the description.** In your editor, replace `(I write this)`. Claude reads only this line when it decides whether to start a skill on its own (*Which Description Gets This Skill Invoked at the Right Moment?*). For comparison, here's the start of a course skill's description. It says *when* it applies, not only what it does:

Coach a participant's context artifact or experiment during the Session 3 exercise. Usage: `/context-coach` — the participant shows their `CLAUDE.local.md` draft ...

Two questions. Which words in your description make Claude reach for this skill right after a commit? Which words would start it at a moment you do not want?

Done when: both files exist, the frontmatter has no `disable-model-invocation` line, and your description says when.

Time left?

Type `/verify-exercise 4`. It grades your two skills now. At the end of Part 2 it grades the hook and `CHANGELOG.md` as well.

Stop here

Part 2 is on a separate sheet. Open it when the trainer starts it, after the slides on hooks. Keep your terminal open: Part 2 uses the same branch and the same skills.